



UCL

Local Deployment of DeepSeek R1 671B



Xihan Li

University College London

Google Developers Expert in Machine Learning

<https://snowkylin.github.io>

DeepSeek R1 671B: one of the first-tier LLMs with reasoning



<https://huggingface.co/deepseek-ai/DeepSeek-R1>
<https://github.com/deepseek-ai/DeepSeek-R1>
<https://arxiv.org/abs/2501.12948>

<https://github.com/lechmazur/writing>


Distilled Model vs Original Model

- Distilled DeepSeek R1 = fine-tuned Qwen (Alibaba) or Llama (Meta) model
 - Use the 671B original model to “teach” other open-sourced models
 - Qwen 1.5B (mobile), 7B (8G+ VRAM), 14B (16G+ VRAM) , 32B (24G+ VRAM)
 - Llama: 8B (8G+ VRAM) / 70B (2 x 24G+ VRAM)
 - Easy deployment (LM Studio)

Model	AIME 2024		MATH-500	GPQA Diamond	LiveCode Bench	CodeForces
	pass@1	cons@64	pass@1	pass@1	pass@1	rating
GPT-4o-0513	9.3	13.4	74.6	49.9	32.9	759
Claude-3.5-Sonnet-1022	16.0	26.7	78.3	65.0	38.9	717
OpenAI-o1-mini	63.6	80.0	90.0	60.0	53.8	1820
QwQ-32B-Preview	50.0	60.0	90.6	54.5	41.9	1316
DeepSeek-R1-Distill-Qwen-1.5B	28.9	52.7	83.9	33.8	16.9	954
DeepSeek-R1-Distill-Qwen-7B	55.5	83.3	92.8	49.1	37.6	1189
DeepSeek-R1-Distill-Qwen-14B	69.7	80.0	93.9	59.1	53.1	1481
DeepSeek-R1-Distill-Qwen-32B	72.6	83.3	94.3	62.1	57.2	1691
DeepSeek-R1-Distill-Llama-8B	50.4	80.0	89.1	49.0	39.6	1205
DeepSeek-R1-Distill-Llama-70B	70.0	86.7	94.5	65.2	57.5	1633
DeepSeek-R1-671B	79.8		97.3	71.5	65.9	2029

Is it possible to deploy a 671B model locally on our own machine?

- 671B means 671,000,000,000 parameters
- ~720GB in its original form
- Typically requires a monster-like AI server like NVIDIA DGX H100
 - 8 x H100 80G
 - ~\$200,000 USD
 - 130.45kg
 - 10.2kW max
- Seems far from customer-grade hardware



<https://www.nvidia.com/en-gb/data-center/dgx-h100/>
https://x.com/netris_io/status/1770593406395515207

Dynamic Quantization

- DeepSeek R1 is trained on FP8
- Standard Quantization: quantize all parameters from 8-bit to x-bit
 - Standard 4-bit quantized version is still 404GB
- Dynamic Quantization:
 - Selectively quantize a few important layers to higher bits (4-6 bits)
 - Leave most of the MoE layers (~88% parameters) to be quantized to lower bits (1.5-2.5 bits)


MoE Bits	Disk Size	Type	Quality	Link	Down_proj
1.58-bit	131GB	IQ1_S	Fair	Link	2.06/1.56bit
1.73-bit	158GB	IQ1_M	Good	Link	2.06bit
2.22-bit	183GB	IQ2_XXS	Better	Link	2.5/2.06bit
2.51-bit	212GB	Q2_K_XL	Best	Link	3.5/2.5bit

Model Size	Dynamic Quant	Model Size	Basic Quant
131GB	6.92	133GB	0
158GB	9.08	149GB	1.67
183GB	9.17	175GB	6.17

<https://unsloth.ai/blog/deepseekr1-dynamic>



Many “impossible” missions can be accomplished if we can do some trade-offs...

- Precision
 - 8-bit (720GB) – 4-bit (405GB) – **2.51~1.53-bit (212GB-131GB)**
- Target speed / context window
 - 50-100 token/s, max 128k tokens – **2-10 token/s, max 4096 tokens**
- Budget
 - \$500-3000 (typical customer-grade PC) – **\$5600+ (high-end PC)**



Is it worth it?


Cloud Service (Web/API)



```
# Please install OpenAI SDK first: 'pip3 install openai'
from openai import OpenAI
client = OpenAI(api_key="<DeepSeek API Key>", base_url=
response = client.chat.completions.create(
    model="deepseek-chat",
    messages=[
        {"role": "system", "content": "You are a helpful"},
        {"role": "user", "content": "Hello"},
    ],
    stream=False
)
print(response.choices[0].message.content)
```

- Cheap
- Convenient
- Fast
- Larger context window
- (should be) more stable
- Less restrictive when using third-party hosted model (e.g., in Cursor)

Local Deployment



- Fully offline
 - Run without stable Internet (robot, car, embedded/industrial devices, etc.)
 - Less privacy concerns
- Possibility of full customization
 - Less restrictive, can be fine-tuned
- Not very costly if existing hardware can be utilized
- It is so cool, isn't it?

Hardware Requirement

- Memory (RAM + VRAM) size and bandwidth is the main constraint
 - 1.73-bit dynamic quantization: RAM + VRAM \geq 200 GB
 - 4-bit standard quantization: RAM + VRAM \geq 500 GB
- Mac Studio: 192GB unified memory (\$5600)
- Maximal memory that can be installed on a 4-slot customer-grade motherboard: 4x48GB = 192GB (VENGEANCE® RGB 192GB (4x48GB) DDR5 DRAM 5200MT/s, ~\$750)
 - Buy a kit with four memory sticks for stability!
- Maximal memory that can be installed on a 4-slot workstation motherboard (TRX50 for ThreadRipper): 4x96GB = 384GB (Micron 96GB 5600MHz ECC Registered 2Rx4 DDR5 Server Memory, ~\$2400)
 - 128GB is possible but super expensive
 - 8-slot workstation motherboard (WRX90) is also available but expensive
- VRAM (GPU) is optional
 - Offload some layers to VRAM for acceleration

Deploy with Ollama

- Popular choice: llama.cpp, ollama, vLLM
 - Llama.cpp requires manual compilation for CUDA support (failed for me)
 - If you don't use GPU, llama.cpp is a good choice, see <https://unsloth.ai/blog/deepseekr1-dynamic> for instruction
 - vLLM's GGUF format support is not good
 - Required by the dynamic quantization version
 - If you don't use GGUF, vLLM is a good choice

Deployment process

1. Download .gguf files from <https://huggingface.co/unsloth/DeepSeek-R1-GGUF> and merge with llama.cpp's "llama-gguf-split" command

2. Install Ollama

```
curl -fsSL https://ollama.com/install.sh | sh
```

3. Create a `modelfile` that guide ollama to create a model

4. Create the model in ollama

```
ollama create DeepSeek-R1-UD-IQ1_M -f modelfile
```

5. Run the model

```
ollama run DeepSeek-R1-UD-IQ1_M
```

```
FROM /home/snowkylin/DeepSeek-R1-UD-IQ1_M.gguf
PARAMETER num_gpu XXX
PARAMETER num_ctx 2048
PARAMETER temperature 0.6
TEMPLATE "< | User | >{{ .System }} {{ .Prompt }}< | Assistant | >"
```

Thank You!

Xihan Li
University College London
Google Developers Expert in Machine Learning
<https://snowkylin.github.io>